

CSE471: Statistical Methods in AI

Assignment 3

Abhinav Moudgil [201331039]

Contents

- Naive Bayes classifier for discrete data
- BPE derivation for univariate normal density function
- BPE derivation for multivariate normal density function
- Gaussian Naive Bayes classifier with PCA and LDA
- References

Naive Bayes classifier for discrete data

Dataset: UCI Bank Marketing Data Set¹

The data is related with direct marketing campaigns of a Portuguese banking institution. The marketing campaigns were based on phone calls. Often, more than one contact to the same client was required, in order to access if the product (bank term deposit) would be ('yes') or not ('no') subscribed.

Pre processing:

Original data set contained multivariate data. It was trimmed to contain only discrete and categorical details. Following are the details of data set after trimming:

Characteristics: discrete, categorical

Number of instances: 41188

Number of features: 14

Classifier:

Naive Bayes classifier uses the concept of probability to classify new entities. It assumes that all the features are independent of each other. We can thus uncouple each feature and treat each one as independent. That's why it is called *naive* Bayes. We can also call each feature are evidence.

Given multiple evidences or set of features, we have to predict its class from given set of classes.

Thus,

$$P(\text{outcome}/\text{multiple evidence}) = P(\text{evidence1}/\text{outcome}).P(\text{evidence2}/\text{outcome}) \\ \dots P(\text{evidenceN}/\text{outcome}).P(\text{multiple evidence})$$

Or,

$$P(\text{outcome}/\text{evidence}) \text{ or } P(\text{posterior}) = P(\text{likelihood of evidence}).\text{Prior of evidence}/P(\text{evidence})$$

Naive Bayes decision rule:

Sample belongs to the class which has highest posterior probability.

Tie-break: If posterior probabilities for the classes are equal, sample is assigned the class which has highest prior probability because the class with more prior probability means it came more often in history. We assume that the fashion will follow for future, and hence the decision.

Handling of missing data points:

1. Unknown values are treated as another categorical values. Probability of this value is calculated for each feature and used in calculation of test sample probability. Underlying assumption is that number of unknown values in each feature in testing set will follow the same fashion as in training set.
2. There are some values for each feature, for which probability is not calculated while training as it was not there in training set. In such case, that feature is ignored and other features are used for classification.

Code:

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 from sklearn import preprocessing
4 import random
5 import copy
6 import math
7
8 def encode(vector):
9     le = preprocessing.LabelEncoder()
10    le.fit(vector)
11    vector = le.transform(vector)
12    return vector
13
14 def encodeAll(mat):
15     categoricalIndices = [1,2,3,4,5,6,7,8,9,14]
16     for i in categoricalIndices:
17         mat[:,i] = encode(mat[:,i])
18     return mat.astype(np.float)
19
20 def deleteFloatData(mat):
21     idx_OUT_columns = [10,15,16,17,18,19]
22     idx_IN_columns = [i for i in xrange(np.shape(mat)[1]) if i not
23                        in idx_OUT_columns]
24     extractedData = mat[:,idx_IN_columns]
25     return extractedData
26
27 def findProbabilityMatrix0(featureMatrix, C0):
28     l = C0.shape[1]
29     p0 = copy.deepcopy(featureMatrix)
30     for i in xrange(l-1):
31         featureVector = featureMatrix[i]
32         x0 = C0[:, i].tolist()
```

```

32         j = 0
33         for item in featureVector:
34             c0 = float(x0.count(item))
35             p0[i][j] = c0/float(C0.shape[0])
36             j = j + 1
37     return p0
38
39 def findProbabilityMatrix1(featureMatrix, C1):
40     l = C1.shape[1]
41     p1 = copy.deepcopy(featureMatrix)
42     for i in xrange(0, l - 1):
43         featureVector = featureMatrix[i]
44         x1 = C1[:, i].tolist()
45         j = 0
46         for item in featureVector:
47             c1 = float(x1.count(item))
48             p1[i][j] = c1/float(C1.shape[0])
49             j = j + 1
50     return p1
51
52 def findFeatureMatrix(mat):
53     l = mat.shape[1]
54     featureMatrix = []
55     for i in xrange(0, l - 1):
56         featureVector = np.unique(mat[:, i])
57         featureMatrix.append(featureVector.tolist())
58     return featureMatrix
59
60 #Training phase
61 file = open('bank-additional.csv')
62 featureVectors = []
63 i = 0
64 for line in file:
65     vector = line.strip().lower().split(';')
66     if i != 0:
67         if vector[-1] == "no":
68             vector[-1] = 0
69         else:
70             vector[-1] = 1
71     featureVectors.append(vector)
72     i = i + 1
73
74 numberOfRuns = 10
75 results = np.zeros([numberOfRuns])
76 for t in xrange(numberOfRuns):
77     random.shuffle(featureVectors)
78     mat = np.array(featureVectors)
79     mat = encodeAll(mat)
80     mat = deleteFloatData(mat)
81     mat = mat.astype(int)
82     N = 2500
83     trainData = mat[:N, :]
84     testData = mat[N:, :]
85     featureMatrix = findFeatureMatrix(trainData)
86     C0 = trainData[trainData[:, -1] == 0]
87     C1 = trainData[trainData[:, -1] == 1]
88     p0 = findProbabilityMatrix0(featureMatrix, C0)

```

```

89     p1 = findProbabilityMatrix1(featureMatrix, C1)
90     pr0 = float(C0.shape[0])/float(trainData.shape[0])
91     pr1 = float(C1.shape[0])/float(trainData.shape[0])
92     #print pr0
93     #print pr1
94     if (pr0 > pr1):
95         maxPrior = int(0)
96     else:
97         maxPrior = int(1)
98     #Testing phase
99     totalValues = testData.shape[0]
100    #print totalValues
101    #print testData[:, 10]
102    myPrediction = np.zeros([totalValues])
103    for i in xrange(0, totalValues):
104        sample = testData[i, :]
105        sample = sample.tolist()
106        ans0 = (float(pr0))
107        ans1 = (float(pr1))
108        count = 0;
109        for j in xrange(0, len(sample) - 1):
110            flag = 0
111            for k in xrange(0, len(featureMatrix[j])):
112                if (sample[j] == featureMatrix[j][k]):
113                    if (p0[j][k] != 0):
114                        ans0 = ans0 * (p0[j][k])
115                    else:
116                        ans0 = ans0 * p0[j][k]
117                    if (p1[j][k] != 0):
118                        ans1 = ans1 * (p1[j][k])
119                    else:
120                        ans1 = ans1 * p1[j][k]
121                # print "Found!"
122                count = count + 1
123                flag = 1
124                break
125            #if (flag == 0):
126            #    print sample[j], " ", i + trainData.shape[0], " -
127            #    print ans0
128            #    print ans1
129            if (ans0 > ans1):
130                myPrediction[i] = int(0)
131            elif (ans0 < ans1):
132                myPrediction[i] = int(1)
133            else:
134                myPrediction[i] = maxPrior
135
136    trueAns = testData[:, -1]
137    #print len(myPrediction[myPrediction == 0])
138    #print len(myPrediction[myPrediction == 1])
139    #print len(trueAns[trueAns == 0])
140    #print len(trueAns[trueAns == 1])
141    correctValues = 0
142    for i in range(totalValues):
143        if (myPrediction[i] == trueAns[i]):
144            correctValues = correctValues + 1

```


```

145     #else:
146     #     print i + trainData.shape[0]
147
148     correctValues = float(correctValues)
149     totalValues = float(totalValues)
150     accuracy = correctValues/totalValues * 100
151     #print correctValues
152     results[t] = accuracy
153     #print accuracy
154
155 meanAccuracy = np.mean(results)
156 x = np.array([1,2,3,4,5,6,7,8,9,10])
157 y = results
158 plt.figure()
159 #plt.plot(x,y, marker='o', color='b')
160 plt.errorbar(x, y, np.std(results, axis = 0))
161 plt.show()
162 #stD = np.std(results)
163 print meanAccuracy

```

Results:

Size of data set	Mean Accuracy	Standard deviation	Time(s)
10%	88.472	0.311	2.630
100%	88.882	0.352	32.244



Observations:

1. Time complexity of this classifier during testing is $O(MN)$, where N is the number of test samples, and M is the number of features.
2. Mean accuracy of classifier can be increased by taking more training samples.

We'll get more unique feature values for which probability can be calculated. Hence less features will be ignored.

3. It assumes that all the features are independent is not true in actual and thus accuracy is less than 90%.

Bayesian parameter estimation derivation

As done in MLE first will start with a simple case with only the mean: μ unknown. As usual we will assume sample x_k is normally distributed as:

$$p(x|\mu) \sim N(\mu, \sigma^2)$$

and the parameter μ has the distribution of:

$$p(\mu) \sim N(\mu_0, \sigma_0^2),$$

as parameter μ is not estimated to be a number but a random variable.

Using Bayes' formula and the corresponding derivation from the previous section the corresponding function could be easily obtained:

$$p(\mu|D) = \alpha \prod_{k=1}^n p(x_k|\mu)p(\mu)$$

Where α is introduced as a 'scale' coefficient in order to simplify the derivation. Please note that α is completely independent of μ .

As x_k is normally distributed we update the $p(x_k|\mu)$ and $p(u)$ with the known distribution function:

$$p(x_k|\mu) = \frac{1}{(2\pi\sigma^2)^{1/2}} \exp\left[-\frac{1}{2}\left(\frac{x_k - \mu}{\sigma}\right)^2\right]$$

$$p(u) = \frac{1}{(2\pi\sigma_0^2)^{1/2}} \exp\left[-\frac{1}{2}\left(\frac{\mu - \mu_0}{\sigma_0}\right)^2\right]$$

Again, substitute $p(x_k|\mu)$ and $p(u)$ in equation $p(\mu|D) = \alpha \prod_{k=1}^n p(x_k|\mu)p(\mu)$ we obtained:

$$p(\mu|D) = \alpha \prod_{k=1}^n \frac{1}{(2\pi\sigma^2)^{1/2}} \exp\left[-\frac{1}{2}\left(\frac{x_k - \mu}{\sigma}\right)^2\right] \frac{1}{(2\pi\sigma_0^2)^{1/2}} \exp\left[-\frac{1}{2}\left(\frac{\mu - \mu_0}{\sigma_0}\right)^2\right]$$

$$p(\mu|D) = \alpha \prod_{k=1}^n \frac{1}{(2\pi\sigma^2)^{1/2}} \frac{1}{(2\pi\sigma_0^2)^{1/2}} \exp\left[-\frac{1}{2}\left(\frac{\mu - \mu_0}{\sigma_0}\right)^2 - \frac{1}{2}\left(\frac{x_k - \mu}{\sigma}\right)^2\right]$$

Similarly, in order to simplify the derivation we update the scaling factors to α' and α'' , and correspondingly,

$$p(\mu|D) = \alpha' \exp \sum_{k=1}^n \left(-\frac{1}{2}\left(\frac{\mu - \mu_0}{\sigma_0}\right)^2 - \frac{1}{2}\left(\frac{x_k - \mu}{\sigma}\right)^2 \right)$$

$$p(\mu|D) = \alpha'' \exp[-\frac{1}{2}(\frac{n}{\sigma^2} + \frac{1}{\sigma_0^2})\mu^2 - 2(\frac{1}{\sigma^2} \sum_{k=1}^n x_k + \frac{\mu_0}{\sigma_0^2})\mu]$$

Finally, compare derived $p(u|D)$ to the Gaussian Distribution in the standard form:

$$p(u|D) = \frac{1}{(2\pi\sigma_n^2)^{1/2}} \exp[-\frac{1}{2}(\frac{\mu - \mu_n}{\sigma_n})^2]$$

Based on knowledge on Gaussian Distribution, μ_n and σ_n^2 could be obtained accordingly:

$$\begin{aligned} \mu_n &= (\frac{n\sigma_0^2}{n\sigma_0^2 + \sigma^2})\bar{x}_n + \frac{\sigma^2}{n\sigma_0^2 + \sigma^2}\mu_0 \\ \sigma_n^2 &= \frac{\sigma_0^2\sigma^2}{n\sigma_0^2 + \sigma^2} \end{aligned}$$

Univariate case:

Given the posteriori density $p(\mu|D)$ successfully derived (variance: σ_n^2 and mean: μ_n now known), the final step is to estimate $p(x|D)$ based on the conclusions above.

$$\begin{aligned} p(x|\mathcal{D}) &= \int p(x|\mu)p(\mu|\mathcal{D})d\mu \\ p(x|\mathcal{D}) &= \int \frac{1}{\sqrt{2\pi}\sigma} \exp[-\frac{1}{2}(\frac{x - \mu}{\sigma})^2] \frac{1}{\sqrt{2\pi}\sigma_n} \exp[-\frac{1}{2}(\frac{\mu - \mu_n}{\sigma_n})^2]d\mu \end{aligned}$$

Finally, substitute σ_n^2 and μ_n the probability function $p(x|\mathcal{D})$ is obtained:

$$p(x|\mathcal{D}) = \frac{1}{2\pi\sigma\sigma_n} \exp[-\frac{1}{2}\frac{(x - \mu)}{\sigma^2 + \sigma_n^2}] \int \exp[-\frac{1}{2}\frac{\sigma^2 + \sigma_n^2}{\sigma^2\sigma_n^2}(\mu - \frac{\sigma_n^2\bar{x}_n + \sigma^2\mu_n}{\sigma^2 + \sigma_n^2})^2]d\mu$$

Hence, $p(x|D)$ is normally distributed as:

$$p(x|D) \sim N(\mu_n, \sigma^2 + \sigma_n^2)$$

Multivariate case:

The treatment of the multivariate case in which \sum is known but μ is not, is a direct generalization of the univariate case. For this reason we shall only sketch the derivation. As before, we assume that,

$p(x|\sum) \sim N(\mu, \sum)$ and $p(\mu) \sim N(\mu_0, \sum_0)$ We know,

$$p(\mu|D) = \alpha \prod_{k=1}^n p(x_k|\mu)p(\mu),$$

Equating coefficients and simplified by knowledge of the matrix identity,

$$(A^{-1} + B^{-1})^{-1} = A(A + B)^{-1}B = B(A + B)^{-1}A,$$

We obtain the final results,

$$\mu_n = \Sigma_0(\Sigma_0 + \frac{1}{n}\Sigma)^{-1}(\frac{1}{n}\sum_{i=1}^n x_i) + \frac{1}{n}\Sigma(\Sigma_0 + \frac{1}{n}\Sigma)^{-1}\mu_0$$

and,

$$\Sigma_n = \Sigma_0(\Sigma_0 + \frac{1}{n}\Sigma)^{-1}\frac{1}{n}\Sigma$$

Gaussian Naive Bayes classifier with PCA and LDA

Dataset: Arcene²

It was obtained by merging three mass-spectrometry datasets to obtain enough training and test data for a benchmark. The original features indicate the abundance of proteins in human sera having a given mass value. Based on those features one must separate cancer patients from healthy patients. We added a number of distractor feature called *probes* having no predictive power. The order of the features and patterns were randomized.

Attribute characteristics: Real

Number of Instances: 900

Number of features:

Real: 7000

Probes: 3000

Total: 10000

Pre processing:

Data set originally contained 10000 features. Number of features are reduced by two approaches:

Principal Component Analysis (PCA).

PCA reduces the dimensions without taking into account the separation of classes. It picks up the top k dimensions with maximum variance, which facilitates separation of classes. The reduction procedure is explained below:

1. Covariance matrix is found for different features. It gives the idea of behaviour of change of one dimension with respect to another.
2. Eigen values of covariance matrix are found which gives the variance across different dimensions. They are sorted in descending order and top k dimensions are picked.
3. New data is found by:

$$NewData = RowFeatureVector.RowDataAdjust$$

where *RowFeatureVector* is the matrix with the eigen vectors in the columns transposed so that the eigen vectors are now in the rows, with the most significant eigen vector at the top, and *RowDataAdjust* is the mean-adjusted data transposed, ie. the data items are in each column, with each row holding a separate dimension.

Linear Discriminant Analysis (LDA)

LDA maximizes between class separation i.e variance along resultant dimension is maximized.

It is done in 5 steps as below:

1. First, mean vector for each class are found.
2. Within class scatter matrix SW and between class scatter matrix SB is found.

$$SW = \sum_{i=1}^c \sum_{x=1}^N (x - m_i)(x - m_i)^T$$

$$SB = \sum_{i=1}^c N_i(m_i - m)(m_i - m)^T$$

3. Calculate the eigen values and eigen vectors of matrix $SW^{-1}SB$
4. Eigen vector are sorted by decreasing eigen values and first one is picked.
5. New data is found by:

$$NewData = RowFeatureVector.RowDataAdjust$$

Classifier:

Same bayesian decision rule is followed i.e class with highest posterior probability is assigned to test the sample. In *gaussian* naive bayes classifier, it is assumed that probability of each feature follows the gaussian distribution. In training phase, mean and variance are found for each feature. During testing, sample is picked and probability of each feature is found using the following gaussian function:

$$G(\mu, \sigma) = e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

where μ is mean and σ is variance of the feature.

Results:

PCA

K	Mean Accuracy	Time(m)
10	56.0	16.39
100	56.0	16.39
1000	56.0	17.01

LDA

Mean Accuracy	Time(m)
56.0	24.0

Observations:

1. Mean accuracy is low because of random probes added in data which don't play any role in classifications and reduction of features by PCA and LDA.
2. In gaussian classifier, we have assumed that features will follow the gaussian distribution which is not true in real.
3. Execution time is nearly the same because in both the reduction method, we have calculated all the eigen vectors and picked up top k. The only difference lies in the calculation of new data from old one.
4. Sometimes gaussian function returns the value which is very low ($< 10^{-323}$) and python truncates the value to 0. So we can't compare the values, which again results in low accuracy of classifier.
5. LDA calculates the inverse of within class scatter matrix SW . For this dataset, within class scatter matrix SW is singular. Hence pseudo inverse for scatter matrix is calculated, which results in high execution time than PCA.

PCA

```
1 import numpy as np
2 from numpy import linalg as LA
3 import math
4
5 def gaussian(x, v, M, V):
6     G = (math.exp(-(math.pow((v - M[x]), 2)/(2 * V[x]))))
7     return G
8
9 def findVariance(C):
10     return np.var(C, axis = 0)
11
12 def findMean(C):
13     return np.mean(C, axis = 0)
14
15 def pcaTransform(data, k):
16     data = data - np.mean(data, axis = 0)
17     covarianceMatrix = np.cov(np.transpose(data))
18     eigenValues, eigenVectors = LA.eig(covarianceMatrix)
19     s = np.argsort(eigenValues)[::-1]
20     ev = np.zeros(eigenVectors.shape)
21     for i in xrange(eigenValues.shape[0]):
22         ev[:, i] = eigenVectors[:, s[i]]
23     data = np.matrix(data)
24     data = np.transpose(data)
25     eigenVectors = ev[:, :k]
26     eigenVectors = np.matrix(eigenVectors)
27     eigenVectors = np.transpose(eigenVectors)
28     return eigenVectors
```

```

29
30 def addLabels(data, trainLabels):
31     b = np.zeros((data.shape[0], data.shape[1] + 1))
32     b[:, :-1] = data
33     b[:, -1] = trainLabels
34     return b
35
36 def mergeData(trainData, testData):
37     x = np.zeros((trainData.shape[0] + testData.shape[0], trainData
38     .shape[1]))
39     x[:trainData.shape[0], :] = trainData
40     x[trainData.shape[0]:, :] = testData
41     return x
42
43 def project(data, eigenVectors):
44     data = data - np.mean(data, axis = 0)
45     data = np.matrix(data)
46     data = np.transpose(data)
47     newData = eigenVectors * data
48     newData = np.transpose(newData)
49     newData = np.array(newData)
50     return newData
51
52 def getDataMatrix(file, intOrFloat):
53     #intOrFloat decides whether data should be int or float
54     if (intOrFloat == 1):
55         featureVectors = []
56         for line in file:
57             vector = line.strip().lower().split(' ')
58             featureVectors.append(vector)
59         data = np.array(featureVectors)
60         data = data.astype(float)
61     else:
62         trainLabels = []
63         for line in file:
64             vector = line
65             trainLabels.append(vector)
66         data = np.array(trainLabels)
67         data = data.astype(int)
68     return data
69
70 file = open('arcene_train.data.txt')
71 data = getDataMatrix(file, 1)
72 file = open('arcene_train.labels.txt')
73 trainLabels = getDataMatrix(file, 0)
74 file = open('arcene_valid.data.txt')
75 testData = getDataMatrix(file, 1)
76 file = open('arcene_valid.labels.txt')
77 testLabels = getDataMatrix(file, 0)
78 #PCA
79 k = 1000
80 ev = pcaTransform(data, k)
81 trainData = project(data, ev)
82 testData = project(testData, ev)
83 trainData = addLabels(trainData, trainLabels)
84 testData = addLabels(testData, testLabels)

```

```

85 C0 = trainData[trainData[:, -1] == -1]
86 C1 = trainData[trainData[:, -1] == 1]
87 V0 = findVariance(C0[:, :-1])
88 V1 = findVariance(C1[:, :-1])
89 M0 = findMean(C0[:, :-1])
90 M1 = findMean(C1[:, :-1])
91 pr0 = float(C0.shape[0])/float(trainData.shape[0])
92 pr1 = float(C1.shape[0])/float(trainData.shape[0])
93 if (pr0 > pr1):
94     maxPrior = int(-1)
95 else:
96     maxPrior = int(1)
97 print maxPrior
98 L = math.pow(10, -323)
99 MAX = -math.pow(10, 300)
100
101 #Testing phase
102 totalValues = testData.shape[0]
103 myPrediction = np.zeros([totalValues])
104 for i in xrange(0, totalValues):
105     sample = testData[i, :]
106     sample = sample.tolist()
107     ans0 = math.log(float(pr0))
108     ans1 = math.log(float(pr1))
109     count = 0;
110     for j in xrange(0, len(sample) - 1):
111         g1 = gaussian(j, sample[j], M0, V0)
112         g2 = gaussian(j, sample[j], M1, V1)
113         if (g1 < L):
114             ans0 = MAX
115         if (g2 < L):
116             ans1 = MAX
117         if (ans0 > MAX):
118             ans0 = ans0 + math.log(g1)
119         if (ans1 > MAX):
120             ans1 = ans1 + math.log(g2)
121     print "ans0 ", ans0
122     print "ans1 ", ans1
123     if (ans0 > ans1):
124         myPrediction[i] = int(-1)
125     elif (ans1 > ans0):
126         myPrediction[i] = int(1)
127     elif (ans1 == ans0):
128         print "Max - ", maxPrior
129         myPrediction[i] = maxPrior
130 trueAns = testData[:, -1]
131 correctValues = 0
132 for i in range(totalValues):
133     if (myPrediction[i] == trueAns[i]):
134         correctValues = correctValues + 1
135
136 correctValues = float(correctValues)
137 totalValues = float(totalValues)
138 accuracy = correctValues/totalValues * 100
139 print accuracy
140

```

LDA

```

1 import numpy as np
2 from numpy import linalg as LA
3 import math
4 from sklearn import preprocessing
5
6 def gaussian(v, M, V):
7     G = (math.exp(-(math.pow((v - M), 2) / (2 * V))))
8     return G
9
10 def findVariance(C):
11     return np.var(C, axis = 0)
12
13 def findMean(C):
14     return np.mean(C, axis = 0)
15
16 def mergeData(trainData, testData):
17     x = np.zeros((trainData.shape[0] + testData.shape[0], trainData
18     .shape[1]))
19     x[:trainData.shape[0], :] = trainData
20     x[trainData.shape[0]:, :] = testData
21     return x
22
23 def ldaTransform(data):
24     C0 = data[data[:, -1] == -1]
25     C1 = data[data[:, -1] == 1]
26     C0 = C0[:, :-1]
27     C1 = C1[:, :-1]
28     S0 = np.cov(np.transpose(C0))
29     S1 = np.cov(np.transpose(C1))
30     SW = S0 + S1
31     Mu0 = np.mean(C0, axis = 0)
32     Mu1 = np.mean(C1, axis = 0)
33     Mu = np.mean(data, axis = 0)
34     Mu = Mu[:-1]
35     Mu = np.matrix(Mu)
36     Mu0 = np.matrix(Mu0)
37     Mu1 = np.matrix(Mu1)
38     SB = C0.shape[0] * np.transpose(Mu0 - Mu) * (Mu0 - Mu) + C1.
39     shape[0] * np.transpose(Mu1 - Mu) * (Mu1 - Mu)
40     #t = Mu0 - Mu1
41     #t = np.matrix(t)
42     #SB = np.transpose(t) * t
43     Swin = LA.pinv(SW) #costly
44     SwinSB = Swin * SB #costly
45     e, v = LA.eig(SwinSB) #costly
46     s = np.argsort(e)[:, -1]
47     v = np.array(v)
48     ev = np.zeros(v.shape)
49     for i in xrange(e.shape[0]):
50         ev[:, i] = v[:, s[i]]
51     w = ev[:, 0]
52     w = np.matrix(w)
53     return w
54
55 def project(data, w):
56     data = np.matrix(data)

```

```

56     data = np.transpose(data)
57     newData = w * data
58     newData = np.transpose(newData)
59     newData = np.array(newData)
60     return newData
61
62 def addLabels(data, trainLabels):
63     b = np.zeros((data.shape[0], data.shape[1] + 1))
64     b[:, :-1] = data
65     b[:, -1] = trainLabels
66     return b
67
68 def getDataMatrix(file, intOrFloat):
69     #intOrFloat decides whether data should be int or float
70     if (intOrFloat == 1):
71         featureVectors = []
72         for line in file:
73             vector = line.strip().lower().split(' ')
74             featureVectors.append(vector)
75         data = np.array(featureVectors)
76         data = data.astype(float)
77     else:
78         trainLabels = []
79         for line in file:
80             vector = line
81             trainLabels.append(vector)
82         data = np.array(trainLabels)
83         data = data.astype(int)
84     return data
85
86
87 file = open('arcene_train.data.txt')
88 data = getDataMatrix(file, 1)
89 file = open('arcene_train.labels.txt')
90 trainLabels = getDataMatrix(file, 0)
91 file = open('arcene_valid.data.txt')
92 testData = getDataMatrix(file, 1)
93 file = open('arcene_valid.labels.txt')
94 testLabels = getDataMatrix(file, 0)
95
96 #LDA
97 trainData = addLabels(data, trainLabels)
98 ev = ldaTransform(trainData)
99 trainData = trainData[:, :-1]
100 trainData = project(trainData, ev)
101 testData = project(testData, ev)
102 trainData = addLabels(trainData, trainLabels)
103 testData = addLabels(testData, testLabels)
104 C0 = trainData[trainData[:, -1] == -1]
105 C1 = trainData[trainData[:, -1] == 1]
106 V0 = findVariance(C0[:, :-1])
107 V1 = findVariance(C1[:, :-1])
108 M0 = findMean(C0[:, :-1])
109 M1 = findMean(C1[:, :-1])
110 pr0 = float(C0.shape[0])/float(trainData.shape[0])
111 pr1 = float(C1.shape[0])/float(trainData.shape[0])
112 L = math.pow(10, -323)

```

```

113 MAX = -math.pow(10, 300)
114
115 #Testing phase
116 totalValues = testData.shape[0]
117 myPrediction = np.zeros([totalValues])
118 j = 0
119 for i in xrange(0, totalValues):
120     sample = testData[i, :]
121     sample = sample.tolist()
122     ans0 = math.log(float(pr0))
123     ans1 = math.log(float(pr1))
124     count = 0;
125     g1 = gaussian(sample[j], M0, V0)
126     g2 = gaussian(sample[j], M1, V1)
127     if (g1 < L):
128         ans0 = MAX
129     if (g2 < L):
130         ans1 = MAX
131     if (ans0 > MAX):
132         ans0 = ans0 + math.log(g1)
133     if (ans1 > MAX):
134         ans1 = ans1 + math.log(g2)
135     print "ans0 ", ans0
136     print "ans1 ", ans1
137     if (ans0 > ans1):
138         myPrediction[i] = int(-1)
139     else:
140         myPrediction[i] = int(1)
141
142 trueAns = testData[:, -1]
143 correctValues = 0
144 for i in range(totalValues):
145     if (myPrediction[i] == trueAns[i]):
146         correctValues = correctValues + 1
147 correctValues = float(correctValues)
148 totalValues = float(totalValues)
149 accuracy = correctValues/totalValues * 100
150 print accuracy

```

References

1. S. Moro, P. Cortez and P. Rita. A Data-Driven Approach to Predict the Success of Bank Telemarketing. Decision Support Systems, Elsevier, 62:22-31, June 2014
2. Isabelle Guyon, Steve R. Gunn, Asa Ben-Hur, Gideon Dror, 2004. Result analysis of the NIPS 2003 feature selection challenge.