

CSE471: Statistical Methods in AI

Assignment 4

Abhinav Moudgil [201331039]

Contents

- Non linear Fisher's LDA derivation
- Kernel PCA and Kernel LDA
- SVM classifier with kernel PCA and LDA
- References

Non linear Fisher's LDA derivation

In statistics, kernel Fisher discriminant analysis (KFD), also known as generalized discriminant analysis and kernel discriminant analysis, is a kernelized version of linear discriminant analysis. It is named after Ronald Fisher. Using the kernel trick, LDA is implicitly performed in a new feature space, which allows non-linear mappings to be learned.

Linear discriminant analysis

Intuitively, the idea of LDA is to find a projection where class separation is maximized. Given two sets of labeled data, C_1 and C_2 , define the class means \mathbf{m}_1 and \mathbf{m}_2 to be

$$\mathbf{m}_i = \frac{1}{l_i} \sum_{n=1}^{l_i} \mathbf{x}_n^i,$$

where l_i is the number of examples of class C_i . The goal of linear discriminant analysis is to give a large separation of the class means while also keeping the in-class variance small. This is formulated as maximizing

$$J(\mathbf{w}) = \frac{\mathbf{w}^T \mathbf{S}_B \mathbf{w}}{\mathbf{w}^T \mathbf{S}_W \mathbf{w}},$$

where \mathbf{S}_B is the between-class covariance matrix and \mathbf{S}_W is the total within-class covariance matrix:

$$\begin{aligned} \mathbf{S}_B &= (\mathbf{m}_2 - \mathbf{m}_1)(\mathbf{m}_2 - \mathbf{m}_1)^T \\ \mathbf{S}_W &= \sum_{i=1,2} \sum_{n=1}^{l_i} (\mathbf{x}_n^i - \mathbf{m}_i)(\mathbf{x}_n^i - \mathbf{m}_i)^T. \end{aligned}$$

Differentiating $J(\mathbf{w})$ with respect to \mathbf{w} , setting equal to zero, and rearranging gives

$$(\mathbf{w}^T \mathbf{S}_B \mathbf{w}) \mathbf{S}_W \mathbf{w} = (\mathbf{w}^T \mathbf{S}_W \mathbf{w}) \mathbf{S}_B \mathbf{w}.$$

Since we only care about the direction of \mathbf{w} and $\mathbf{S}_B \mathbf{w}$ has the same direction as $(\mathbf{m}_2 - \mathbf{m}_1)$, $\mathbf{S}_B \mathbf{w}$ can be replaced by $(\mathbf{m}_2 - \mathbf{m}_1)$ and we can drop the scalars $(\mathbf{w}^\top \mathbf{S}_B \mathbf{w})$ and $(\mathbf{w}^\top \mathbf{S}_W \mathbf{w})$ to give

$$\mathbf{w} \propto \mathbf{S}_W^{-1} (\mathbf{m}_2 - \mathbf{m}_1).$$

Kernel trick with LDA

To extend LDA to non-linear mappings, the data can be mapped to a new feature space, F , via some function ϕ . In this new feature space, the function that needs to be maximized is

$$J(\mathbf{w}) = \frac{\mathbf{w}^\top \mathbf{S}_B^\phi \mathbf{w}}{\mathbf{w}^\top \mathbf{S}_W^\phi \mathbf{w}},$$

where

$$\begin{aligned} \mathbf{S}_B^\phi &= (\mathbf{m}_2^\phi - \mathbf{m}_1^\phi)(\mathbf{m}_2^\phi - \mathbf{m}_1^\phi)^\top \\ \mathbf{S}_W^\phi &= \sum_{i=1,2} \sum_{n=1}^{l_i} (\phi(\mathbf{x}_n^i) - \mathbf{m}_i^\phi)(\phi(\mathbf{x}_n^i) - \mathbf{m}_i^\phi)^\top, \end{aligned}$$

and

$$\mathbf{m}_i^\phi = \frac{1}{l_i} \sum_{j=1}^{l_i} \phi(\mathbf{x}_j^i).$$

Further, note that $\mathbf{w} \in F$. Explicitly computing the mappings $\phi(\mathbf{x}_i)$ and then performing LDA can be computationally expensive, and in many cases intractable. For example, F may be infinitely dimensional. Thus, rather than explicitly mapping the data to F , the data can be implicitly embedded by rewriting the algorithm in terms of dot products and using the kernel trick in which the dot product in the new feature space is replaced by a kernel function, $k(\mathbf{x}, \mathbf{y}) = \phi(\mathbf{x}) \cdot \phi(\mathbf{y})$. LDA can be reformulated in terms of dot products by first noting that \mathbf{w} will have an expansion of the form[5]

$$\mathbf{w} = \sum_{i=1}^l \alpha_i \phi(\mathbf{x}_i).$$

Then note that

$$\mathbf{w}^\top \mathbf{m}_i^\phi = \frac{1}{l_i} \sum_{j=1}^l \sum_{k=1}^{l_i} \alpha_j k(\mathbf{x}_j, \mathbf{x}_k^i) = \alpha^\top \mathbf{M}_i,$$

where

$$(\mathbf{M}_i)_j = \frac{1}{l_i} \sum_{k=1}^{l_i} k(\mathbf{x}_j, \mathbf{x}_k^i).$$

The numerator of $J(\mathbf{w})$ can then be written as:

$$\begin{aligned}\mathbf{w}^\top \mathbf{S}_B^\phi \mathbf{w} &= \mathbf{w}^\top (\mathbf{m}_2^\phi - \mathbf{m}_1^\phi)(\mathbf{m}_2^\phi - \mathbf{m}_1^\phi)^\top \mathbf{w} \\ &= \alpha^\top \mathbf{M} \alpha,\end{aligned}$$

where $\mathbf{M} = (\mathbf{M}_2 - \mathbf{M}_1)(\mathbf{M}_2 - \mathbf{M}_1)^\top$. Similarly, the denominator can be written as

$$\mathbf{w}^\top \mathbf{S}_W^\phi \mathbf{w} = \alpha^\top \mathbf{N} \alpha,$$

where

$\mathbf{N} = \sum_{j=1,2} \mathbf{K}_j (\mathbf{I} - \mathbf{1}_{l_j}) \mathbf{K}_j^\top$, with the $n^{\text{th}}, m^{\text{th}}$ component of \mathbf{K}_j defined as $k(\mathbf{x}_n, \mathbf{x}_m^j)$, \mathbf{I} is the identity matrix, and $\mathbf{1}_{l_j}$ the matrix with all entries equal to $1/l_j$. This identity can be derived by starting out with the expression for $\mathbf{w}^\top \mathbf{S}_W^\phi \mathbf{w}$ and using the expansion of \mathbf{w} and the definitions of \mathbf{S}_W^ϕ and \mathbf{m}_i^ϕ

$$\begin{aligned}\mathbf{w}^\top \mathbf{S}_W^\phi \mathbf{w} &= \left(\sum_{i=1}^l \alpha_i \phi^\top(\mathbf{x}_i) \right) \left(\sum_{j=1,2} \sum_{n=1}^{l_j} (\phi(\mathbf{x}_n^j) - \mathbf{m}_j^\phi)(\phi(\mathbf{x}_n^j) - \mathbf{m}_j^\phi)^\top \right) \left(\sum_{k=1}^l \alpha_k \phi(\mathbf{x}_k) \right) \\ &= \sum_{j=1,2} \sum_{i=1}^l \sum_{n=1}^{l_j} \sum_{k=1}^l \alpha_i \phi^\top(\mathbf{x}_i) (\phi(\mathbf{x}_n^j) - \mathbf{m}_j^\phi)(\phi(\mathbf{x}_n^j) - \mathbf{m}_j^\phi)^\top \alpha_k \phi(\mathbf{x}_k) \\ &= \sum_{j=1,2} \sum_{i=1}^l \sum_{n=1}^{l_j} \sum_{k=1}^l \left(\alpha_i k(\mathbf{x}_i, \mathbf{x}_n^j) - \frac{1}{l_j} \sum_{p=1}^{l_j} \alpha_i k(\mathbf{x}_i, \mathbf{x}_p^j) \right) \left(\alpha_k k(\mathbf{x}_k, \mathbf{x}_n^j) - \frac{1}{l_j} \sum_{q=1}^{l_j} \alpha_k k(\mathbf{x}_k, \mathbf{x}_q^j) \right) \\ &= \sum_{j=1,2} \left(\sum_{i=1}^l \sum_{n=1}^{l_j} \sum_{k=1}^l \left(\alpha_i \alpha_k k(\mathbf{x}_i, \mathbf{x}_n^j) k(\mathbf{x}_k, \mathbf{x}_n^j) \right. \right. \\ &\quad \left. \left. - \frac{2\alpha_i \alpha_k}{l_j} \sum_{p=1}^{l_j} k(\mathbf{x}_i, \mathbf{x}_p^j) k(\mathbf{x}_k, \mathbf{x}_p^j) + \frac{\alpha_i \alpha_k}{l_j^2} \sum_{p=1}^{l_j} \sum_{q=1}^{l_j} k(\mathbf{x}_i, \mathbf{x}_p^j) k(\mathbf{x}_k, \mathbf{x}_q^j) \right) \right) \\ &= \sum_{j=1,2} \left(\sum_{i=1}^l \sum_{n=1}^{l_j} \sum_{k=1}^l \left(\alpha_i \alpha_k k(\mathbf{x}_i, \mathbf{x}_n^j) k(\mathbf{x}_k, \mathbf{x}_n^j) - \frac{\alpha_i \alpha_k}{l_j} \sum_{p=1}^{l_j} k(\mathbf{x}_i, \mathbf{x}_p^j) k(\mathbf{x}_k, \mathbf{x}_p^j) \right) \right) \\ &= \sum_{j=1,2} \alpha^\top \mathbf{K}_j \mathbf{K}_j^\top \alpha - \alpha^\top \mathbf{K}_j \mathbf{1}_{l_j} \mathbf{K}_j^\top \alpha \\ &= \alpha^\top \mathbf{N} \alpha.\end{aligned}$$

With these equations for the numerator and denominator of $J(\mathbf{w})$, the equation for J can be rewritten as

$$J(\alpha) = \frac{\alpha^\top \mathbf{M} \alpha}{\alpha^\top \mathbf{N} \alpha}.$$

Then, differentiating and setting equal to zero gives

$$(\alpha^T \mathbf{M} \alpha) \mathbf{N} \alpha = (\alpha^T \mathbf{N} \alpha) \mathbf{M} \alpha.$$

Since only the direction of \mathbf{w} , and hence the direction of α , matters, the above can be solved for α as

$$\alpha = \mathbf{N}^{-1}(\mathbf{M}_2 - \mathbf{M}_1).$$

Note that in practice, \mathbf{N} is usually singular and so a multiple of the identity is added to it

$$\mathbf{N}_\epsilon = \mathbf{N} + \epsilon \mathbf{I}.$$

Given the solution for α , the projection of a new data point is given by

$$y(\mathbf{x}) = (\mathbf{w} \cdot \phi(\mathbf{x})) = \sum_{i=1}^l \alpha_i k(\mathbf{x}_i, \mathbf{x}).$$

Kernel PCA and LDA

Kernel Principal Component Analysis

Kernel PCA reduces the dimensions *without* taking into account the separation of classes. It picks up the top k dimensions with maximum variance, which facilitates separation of classes.

The reduction procedure is explained below:

1. Kernel (similarity) matrix is computed. Following are the two kernels matrices computed in this assignment:

RBF Kernel:

$$\kappa(\mathbf{x}_i, \mathbf{x}_j) = \exp \left(-\gamma \|\mathbf{x}_i - \mathbf{x}_j\|_2^2 \right)$$

Linear Kernel:

$$\kappa(\mathbf{x}_i, \mathbf{x}_j) = \langle \mathbf{x}_i, \mathbf{x}_j^T \rangle$$

γ is taken 15.

2. Since it is not guaranteed that the kernel matrix is centered, we can apply the following equation to do so:

$$\mathbf{K}' = \mathbf{K} - \mathbf{1}_N \mathbf{K} - \mathbf{K} \mathbf{1}_N + \mathbf{1}_N \mathbf{K} \mathbf{1}_N$$

where $\mathbf{1}_N$ is a $N \times N$ matrix with all values equal to $1/N$.

3. Eigenvectors of the centered kernel matrix that correspond to the largest eigenvalues are the data points already projected onto the respective principal components.

Linear Discriminant Analysis (LDA)

LDA maximizes between class separation i.e variance along resultant dimension is maximized.

It is done in 5 steps as below:

1. First, mean vector for each class are found.
2. Within class scatter matrix SW and between class scatter matrix SB is found.

$$SW = \sum_{i=1}^c \sum_{x=1}^N (x - m_i)(x - m_i)^T$$

$$SB = \sum_{i=1}^c N_i(m_i - m)(m_i - m)^T$$

3. Calculate the eigen values and eigen vectors of matrix $SW^{-1}SB$
4. Eigen vector are sorted by decreasing eigen values and first one is picked.
5. New data is found by:

$$NewData = RowFeatureVector.RowDataAdjust$$

RBF kernel PCA

```
import numpy as np
from scipy.spatial.distance import pdist, squareform
from scipy import exp
from scipy.linalg import eigh

def mergeData(trainData, testData):
    x = np.zeros((trainData.shape[0] + testData.shape[0], trainData.shape[1]))
    x[:trainData.shape[0], :] = trainData
    x[trainData.shape[0]:, :] = testData
    return x

def getDataMatrix(file, intOrFloat):
    #intOrFloat decides whether data should be int or float
    if (intOrFloat == 1):
        featureVectors = []
        for line in file :
            vector = line.strip().lower().split(' ')
            featureVectors.append(vector)
        data = np.array(featureVectors)
        data = data.astype(float)
    else:
        trainLabels = []
        for line in file :
            vector = line
            trainLabels.append(vector)
        data = np.array(trainLabels)
```

```

data = data.astype(int)
return data

def addLabels(data, trainLabels):
    b = np.zeros((data.shape[0], data.shape[1] + 1))
    b[:, :-1] = data
    b[:, -1] = trainLabels
    return b

def kPCA(X, gamma, k):
    distances = pdist(X, 'sqeuclidean')
    symmetricDistances = squareform(distances)
    K = exp(-gamma * symmetricDistances)
    N = K.shape[0]
    one_N = np.ones((N,N))/N
    normalizedK = K - one_N.dot(K) - K.dot(one_N)
    + one_N.dot(K).dot(one_N)
    eigenValues, eigenVectors = eigh(normalizedK)
    alphas = np.column_stack((eigenVectors[:, -i] for i in range(1, k+1)))
    lambdas = [eigenValues[-i] for i in range(1, k+1)]
    return alphas

file = open('arcene_train.data.txt')
X = getDataMatrix(file, 1)
file = open('arcene_train.labels.txt')
trainLabels = getDataMatrix(file, 0)
file = open('arcene_valid.data.txt')
testData = getDataMatrix(file, 1)
file = open('arcene_valid.labels.txt')
testLabels = getDataMatrix(file, 0)

K = 100
gamma = 15
fullData = mergeData(X, testData)
Data = kPCA(fullData, gamma, K)
trainData = Data[:X.shape[0], :]
testData = Data[X.shape[0]:, :]

Linear kernel PCA

import numpy as np
from scipy.spatial.distance import pdist, squareform
from scipy import exp
from scipy.linalg import eigh

def mergeData(trainData, testData):
    x = np.zeros((trainData.shape[0] + testData.shape[0], trainData.shape[1]))

```

```

x[:trainData.shape[0], :] = trainData
x[trainData.shape[0]:, :] = testData
return x

def getDataMatrix(file, intOrFloat):
    #intOrFloat decides whether data should be int or float
    if (intOrFloat == 1):
        featureVectors = []
        for line in file :
            vector = line.strip().lower().split(' ')
            featureVectors.append(vector)
        data = np.array(featureVectors)
        data = data.astype(float)
    else:
        trainLabels = []
        for line in file :
            vector = line
            trainLabels.append(vector)
        data = np.array(trainLabels)
        data = data.astype(int)
    return data

def addLabels(data, trainLabels):
    b = np.zeros((data.shape[0], data.shape[1] + 1))
    b[:, :-1] = data
    b[:, -1] = trainLabels
    return b

def kPCA(X, k):
    K = X.dot(X.T)
    N = K.shape[0]
    one_N = np.ones((N,N))/N
    normalizedK = K - one_N.dot(K) - K.dot(one_N)
    + one_N.dot(K).dot(one_N)
    eigenValues, eigenVectors = eigh(normalizedK)
    alphas = np.column_stack((eigenVectors[:, -i] for i in range(1, k+1)))
    lambdas = [eigenValues[-i] for i in range(1, k+1)]
    return alphas

file = open('arcene_train.data.txt')
X = getDataMatrix(file, 1)
file = open('arcene_train.labels.txt')
trainLabels = getDataMatrix(file, 0)
file = open('arcene_valid.data.txt')
testData = getDataMatrix(file, 1)
file = open('arcene_valid.labels.txt')

```

```

testLabels = getDataMatrix(file, 0)

K = 100
fullData = mergeData(X, testData)
Data = kPCA(fullData, K)
trainData = Data[:X.shape[0], :]
testData = Data[X.shape[0]:, :]

LDA

import numpy as np
from numpy import linalg as LA
import math
from sklearn import preprocessing

def mergeData(trainData, testData):
    x = np.zeros((trainData.shape[0] + testData.shape[0], trainData.shape[1]))
    x[:trainData.shape[0], :] = trainData
    x[trainData.shape[0]:, :] = testData
    return x

def ldaTransform(data):
    C0 = data[data[:, -1] == -1]
    C1 = data[data[:, -1] == 1]
    C0 = C0[:, :-1]
    C1 = C1[:, :-1]
    S0 = np.cov(np.transpose(C0))
    S1 = np.cov(np.transpose(C1))
    SW = S0 + S1
    Mu0 = np.mean(C0, axis = 0)
    Mu1 = np.mean(C1, axis = 0)
    Mu = np.mean(data, axis = 0)
    Mu = Mu[:-1]
    Mu = np.matrix(Mu)
    Mu0 = np.matrix(Mu0)
    Mu1 = np.matrix(Mu1)
    SB = C0.shape[0] * np.transpose(Mu0 - Mu) * (Mu0 - Mu)
    + C1.shape[0] * np.transpose(Mu1 - Mu) * (Mu1 - Mu)
    Swin = LA.pinv(SW) #costly
    Swin = np.matrix(Swin)
    SwinSB = Swin * SB #costly
    e, v = LA.eig(SwinSB) #costly
    s = np.argsort(e)[::-1]
    v = np.array(v)
    ev = np.zeros(v.shape)
    for i in xrange(e.shape[0]):
        ev[:, i] = v[:, s[i]]

```



```

w = ev[:, 0]
w = np.matrix(w)
l = data[:, -1]
data = data[:, :-1]
data = np.matrix(data)
data = np.transpose(data)
newData = w * data
newData = np.transpose(newData)
newData = np.array(newData)
newData = addLabels(newData, l)
return newData

def addLabels(data, trainLabels):
b = np.zeros((data.shape[0], data.shape[1] + 1))
b[:, :-1] = data
b[:, -1] = trainLabels
return b

def getDataMatrix(file, intOrFloat):
#intOrFloat decides whether data should be int or float
if (intOrFloat == 1):
featureVectors = []
for line in file :
vector = line.strip().lower().split(' ')
featureVectors.append(vector)
data = np.array(featureVectors)
data = data.astype(float)
else:
trainLabels = []
for line in file :
vector = line
trainLabels.append(vector)
data = np.array(trainLabels)
data = data.astype(int)
return data

file = open('arcene_train.data.txt')
data = getDataMatrix(file, 1)
file = open('arcene_train.labels.txt')
trainLabels = getDataMatrix(file, 0)
file = open('arcene_valid.data.txt')
testData = getDataMatrix(file, 1)
file = open('arcene_valid.labels.txt')
testLabels = getDataMatrix(file, 0)

trainData = addLabels(data, trainLabels)

```

```
testData = addLabels(testData, testLabels)
fullData = mergeData(trainData, testData)

Data = ldaTransform(fullData)
trainData = Data[:data.shape[0], :-1]
testData = Data[data.shape[0]:, :]
```

SVM classifier with Kernel PCA and LDA

Dataset: Arcene²

It was obtained by merging three mass-spectrometry datasets to obtain enough training and test data for a benchmark. The original features indicate the abundance of proteins in human sera having a given mass value. Based on those features one must separate cancer patients from healthy patients. We added a number of distractor feature called *probes* having no predictive power. The order of the features and patterns were randomized.

Attribute characteristics: Real

Number of Instances: 900

Number of features:

Real: 7000

Probes: 3000

Total: 10000

Pre processing:

Data set originally contained 10000 features. Number of features are reduced by Kernel PCA.

Kernel Principal Component Analysis

Kernel PCA reduces the dimensions *without* taking into account the separation of classes. It picks up the top k dimensions with maximum variance, which facilitates separation of classes.

The reduction procedure is explained below:

1. Kernel (similarity) matrix is computed. Following are the two kernels matrices computed in this assignment:

RBF Kernel:

$$\kappa(\mathbf{x}_i, \mathbf{x}_j) = \exp\left(-\gamma \|\mathbf{x}_i - \mathbf{x}_j\|_2^2\right)$$

Linear Kernel:

$$\kappa(\mathbf{x}_i, \mathbf{x}_j) = \langle \mathbf{x}_i, \mathbf{x}_j^T \rangle$$

γ is taken 15.

2. Since it is not guaranteed that the kernel matrix is centered, we can apply the following equation to do so:

$$K' = K - \mathbf{1}_N K - K \mathbf{1}_N + \mathbf{1}_N K \mathbf{1}_N$$

where $\mathbf{1}_N$ is a $N \times N$ matrix with all values equal to $1/N$.

3. Eigenvectors of the centered kernel matrix that correspond to the largest eigenvalues are the data points already projected onto the respective principal components.

Linear Discriminant Analysis (LDA)

LDA maximizes between class separation i.e variance along resultant dimension is maximized.

It is done in 5 steps as below:

1. First, mean vector for each class are found.
2. Within class scatter matrix SW and between class scatter matrix SB is found.

$$SW = \sum_{i=1}^c \sum_{x=1}^N (x - m_i)(x - m_i)^T$$

$$SB = \sum_{i=1}^c N_i(m_i - m)(m_i - m)^T$$

3. Calculate the eigen values and eigen vectors of matrix $SW^{-1}SB$
4. Eigen vector are sorted by decreasing eigen values and first one is picked.
5. New data is found by:

$$NewData = RowFeatureVector.RowDataAdjust$$

Classifier: SVM classifier from the standard scikit-learn python library is used with soft margin $C = 1.0$ and kernel being linear.

Results:

PCA with RBF kernel

K	Mean Accuracy	Time(s)
10	56.0	1.239
100	57.0	1.260

PCA with linear kernel

K	Mean Accuracy	Time(s)
10	56.0	1.059
100	62.0	1.260

LDA

Mean Accuracy	Time(m)
100.0	16.0

Observations:

1. Linear kernel PCA computes exactly the same result as standard PCA method

but it is much faster as it does eigen decomposition and doesn't explicitly computes the covariance matrix.

2. SVM classifier with PCA gives poor performance than Bayesian classifier used in previous assignment.

3. Variation of γ in RBF kernel doesn't affect classification performance of SVM.

4. LDA calculates the inverse of within class scatter matrix SW . For this dataset, within class scatter matrix SW is singular. Hence pseudo inverse for scatter matrix is calculated, which results in high execution time than PCA.

SVM with linear kernel PCA

```
import numpy as np
from scipy.spatial.distance import pdist, squareform
from scipy import exp
from scipy.linalg import eigh
from sklearn import svm, preprocessing
from sklearn.metrics import classification_report as cr

def train(X, y):
    clf = svm.SVC(kernel='linear', C = 1.0, max_iter = -1)
    clf.fit(X, y)
    return clf

def predict(model, vector):
    return model.predict(vector)

def classify(model, featureVectors):
    true = 0
    total = 0
    z = []
    for feature in featureVectors:
        if feature[-1] == predict(model, feature[:-1]):
            true += 1
    z = z + predict(model, feature[:-1]).astype(np.int).tolist()
    total += 1
    data = featureVectors[:, -1].flatten()
    data = data.astype(np.int).tolist()
    print z
    print cr(data, z)
    print "Accuracy : ",
    print (true * 100) / total

def mergeData(trainData, testData):
    x = np.zeros((trainData.shape[0] + testData.shape[0], trainData.shape[1]))
    x[:trainData.shape[0], :] = trainData
```

```

x[trainData.shape[0]:, :] = testData
return x

def getDataMatrix(file, intOrFloat):
    #intOrFloat decides whether data should be int or float
    if (intOrFloat == 1):
        featureVectors = []
        for line in file :
            vector = line.strip().lower().split(' ')
            featureVectors.append(vector)
        data = np.array(featureVectors)
        data = data.astype(float)
    else:
        trainLabels = []
        for line in file :
            vector = line
            trainLabels.append(vector)
        data = np.array(trainLabels)
        data = data.astype(int)
    return data

def addLabels(data, trainLabels):
    b = np.zeros((data.shape[0], data.shape[1] + 1))
    b[:, :-1] = data
    b[:, -1] = trainLabels
    return b

def kPCA(X, gamma, k):
    K = X.dot(X.T)
    N = K.shape[0]
    one_N = np.ones((N,N))/N
    normalizedK = K - one_N.dot(K) - K.dot(one_N)
    + one_N.dot(K).dot(one_N)
    eigenValues, eigenVectors = eigh(normalizedK)
    alphas = np.column_stack((eigenVectors[:, -i] for i in range(1,k+1)))
    lambdas = [eigenValues[-i] for i in range(1,k+1)]
    return alphas

file = open('arcene_train.data.txt')
X = getDataMatrix(file, 1)
file = open('arcene_train.labels.txt')
trainLabels = getDataMatrix(file, 0)
file = open('arcene_valid.data.txt')
testData = getDataMatrix(file, 1)
file = open('arcene_valid.labels.txt')
testLabels = getDataMatrix(file, 0)

```

```

K = 100
gamma = 15
fullData = mergeData(X, testData)
Data = kPCA(fullData, gamma, K)
trainData = Data[:X.shape[0], :]
testData = Data[X.shape[0]:, :]
testData = addLabels(testData, testLabels)
model = train(trainData, trainLabels)
# trainData = addLabels(trainData, trainLabels)
classify(model, testData)

```

SVM with RBF kernel PCA

```

import numpy as np
from scipy.spatial.distance import pdist, squareform
from scipy import exp
from scipy.linalg import eigh
from sklearn import svm, preprocessing
from sklearn.metrics import classification_report as cr

def train(X, y):
    clf = svm.SVC(kernel='linear', C = 1.0, max_iter = -1)
    clf.fit(X, y)
    return clf

def predict(model, vector):
    return model.predict(vector)

def classify(model, featureVectors):
    true = 0
    total = 0
    z = []
    for feature in featureVectors:
        if feature[-1] == predict(model, feature[:-1]):
            true += 1
        z = z + predict(model, feature[:-1]).astype(np.int).tolist()
    # if predict(model, feature[:-1]) == 1:
    #     print "Yes!"
    total += 1
    data = featureVectors[:, -1].flatten()
    data = data.astype(np.int).tolist()
    print z
    print cr(data, z)
    print "Accuracy : ",
    print (true * 100) / total

```

```

def mergeData(trainData, testData):
    x = np.zeros((trainData.shape[0] + testData.shape[0], trainData.shape[1]))
    x[:trainData.shape[0], :] = trainData
    x[trainData.shape[0]:, :] = testData
    return x

def getDataMatrix(file, intOrFloat):
    #intOrFloat decides whether data should be int or float
    if (intOrFloat == 1):
        featureVectors = []
        for line in file :
            vector = line.strip().lower().split(' ')
            featureVectors.append(vector)
        data = np.array(featureVectors)
        data = data.astype(float)
    else:
        trainLabels = []
        for line in file :
            vector = line
            trainLabels.append(vector)
        data = np.array(trainLabels)
        data = data.astype(int)
    return data

def addLabels(data, trainLabels):
    b = np.zeros((data.shape[0], data.shape[1] + 1))
    b[:, :-1] = data
    b[:, -1] = trainLabels
    return b

def kPCA(X, gamma, k):
    distances = pdist(X, 'sqeuclidean')
    symmetricDistances = squareform(distances)
    K = exp(-gamma * symmetricDistances)
    N = K.shape[0]
    one_N = np.ones((N,N))/N
    normalizedK = K - one_N.dot(K) - K.dot(one_N)
    + one_N.dot(K).dot(one_N)
    eigenValues, eigenVectors = eigh(normalizedK)
    alphas = np.column_stack((eigenVectors[:, -i] for i in range(1, k+1)))
    lambdas = [eigenValues[-i] for i in range(1, k+1)]
    return alphas

file = open('arcene_train.data.txt')
X = getDataMatrix(file, 1)

```

```

file = open('arcene_train.labels.txt')
trainLabels = getDataMatrix(file, 0)
file = open('arcene_valid.data.txt')
testData = getDataMatrix(file, 1)
file = open('arcene_valid.labels.txt')
testLabels = getDataMatrix(file, 0)

K = 100
gamma = 15
fullData = mergeData(X, testData)
Data = kPCA(fullData, gamma, K)
trainData = Data[:X.shape[0], :]
testData = Data[X.shape[0]:, :]
testData = addLabels(testData, testLabels)
model = train(trainData, trainLabels)
classify(model, testData)

```

SVM with LDA

```

import numpy as np
from numpy import linalg as LA
import math
from sklearn import preprocessing
from sklearn import svm, preprocessing
from sklearn.metrics import classification_report as cr

def mergeData(trainData, testData):
    x = np.zeros((trainData.shape[0] + testData.shape[0], trainData.shape[1]))
    x[:trainData.shape[0], :] = trainData
    x[trainData.shape[0]:, :] = testData
    return x

def ldaTransform(data):
    C0 = data[data[:, -1] == -1]
    C1 = data[data[:, -1] == 1]
    C0 = C0[:, :-1]
    C1 = C1[:, :-1]
    S0 = np.cov(np.transpose(C0))
    S1 = np.cov(np.transpose(C1))
    SW = S0 + S1
    Mu0 = np.mean(C0, axis = 0)
    Mu1 = np.mean(C1, axis = 0)
    Mu = np.mean(data, axis = 0)
    Mu = Mu[:-1]
    Mu = np.matrix(Mu)
    Mu0 = np.matrix(Mu0)

```



```

Mu1 = np.matrix(Mu1)
SB = C0.shape[0] * np.transpose(Mu0 - Mu) * (Mu0 - Mu)
+ C1.shape[0] * np.transpose(Mu1 - Mu) * (Mu1 - Mu)
Swin = LA.pinv(SW) #costly
Swin = np.matrix(Swin)
SwinSB = Swin * SB #costly
e, v = LA.eig(SwinSB) #costly
s = np.argsort(e)[::-1]
v = np.array(v)
ev = np.zeros(v.shape)
for i in xrange(e.shape[0]):
    ev[:, i] = v[:, s[i]]
w = ev[:, 0]
w = np.matrix(w)
l = data[:, -1]
data = data[:, :-1]
data = np.matrix(data)
data = np.transpose(data)
newData = w * data
newData = np.transpose(newData)
newData = np.array(newData)
newData = addLabels(newData, l)
return newData

def addLabels(data, trainLabels):
    b = np.zeros((data.shape[0], data.shape[1] + 1))
    b[:, :-1] = data
    b[:, -1] = trainLabels
    return b

def getDataMatrix(file, intOrFloat):
    #intOrFloat decides whether data should be int or float
    if (intOrFloat == 1):
        featureVectors = []
        for line in file :
            vector = line.strip().lower().split(' ')
            featureVectors.append(vector)
        data = np.array(featureVectors)
        data = data.astype(float)
    else:
        trainLabels = []
        for line in file :
            vector = line
            trainLabels.append(vector)
        data = np.array(trainLabels)
        data = data.astype(int)

```

```

return data

def train(X, y):
    clf = svm.SVC(kernel='linear', C = 1.0, max_iter = -1)
    clf.fit(X, y)
    return clf

def predict(model, vector):
    return model.predict(vector)

def classify(model, featureVectors):
    true = 0
    total = 0
    z = []
    for feature in featureVectors:
        if feature[-1] == predict(model, feature[:-1]):
            true += 1
    z = z + predict(model, feature[:-1]).astype(np.int).tolist()
    total += 1
    data = featureVectors[:, -1].flatten()
    data = data.astype(np.int).tolist()
    print z
    print cr(data, z)
    print "Accuracy : ",
    print (true * 100) / total

file = open('arcene_train.data.txt')
data = getDataMatrix(file, 1)
file = open('arcene_train.labels.txt')
trainLabels = getDataMatrix(file, 0)
file = open('arcene_valid.data.txt')
testData = getDataMatrix(file, 1)
file = open('arcene_valid.labels.txt')
testLabels = getDataMatrix(file, 0)
#LDA
trainData = addLabels(data, trainLabels)
testData = addLabels(testData, testLabels)
fullData = mergeData(trainData, testData)
Data = ldaTransform(fullData)
trainData = Data[:data.shape[0], :-1]
testData = Data[data.shape[0]:, :]
#testData = ldaTransform(testData)
model = train(trainData, trainLabels)
classify(model, testData)

```

References

1. S. Moro, P. Cortez and P. Rita. A Data-Driven Approach to Predict the Success of Bank Telemarketing. *Decision Support Systems*, Elsevier, 62:22-31, June 2014
2. Isabelle Guyon, Steve R. Gunn, Asa Ben-Hur, Gideon Dror, 2004. Result analysis of the NIPS 2003 feature selection challenge.